

Effect Size, Statistical Power, and Sample Size Requirements for the Bootstrap Likelihood Ratio Test in Latent Class Analysis

John J. Dziak, Stephanie T. Lanza & Xianming Tan

To cite this article: John J. Dziak, Stephanie T. Lanza & Xianming Tan (2014) Effect Size, Statistical Power, and Sample Size Requirements for the Bootstrap Likelihood Ratio Test in Latent Class Analysis, Structural Equation Modeling: A Multidisciplinary Journal, 21:4, 534-552, DOI: [10.1080/10705511.2014.919819](https://doi.org/10.1080/10705511.2014.919819)

To link to this article: <https://doi.org/10.1080/10705511.2014.919819>



Published online: 23 Jul 2014.



Submit your article to this journal [↗](#)



Article views: 3232



View related articles [↗](#)



View Crossmark data [↗](#)



Citing articles: 95 View citing articles [↗](#)

Effect Size, Statistical Power, and Sample Size Requirements for the Bootstrap Likelihood Ratio Test in Latent Class Analysis

John J. Dziak,¹ Stephanie T. Lanza,¹ and Xianming Tan²

¹*The Pennsylvania State University*

²*McGill University*

Selecting the number of different classes that will be assumed to exist in the population is an important step in latent class analysis (LCA). The bootstrap likelihood ratio test (BLRT) provides a data-driven way to evaluate the relative adequacy of a $(K - 1)$ -class model compared to a K -class model. However, very little is known about how to predict the power or the required sample size for the BLRT in LCA. Based on extensive Monte Carlo simulations, we provide practical effect size measures and power curves that can be used to predict power for the BLRT in LCA given a proposed sample size and a set of hypothesized population parameters. Estimated power curves and tables provide guidance for researchers wishing to size a study to have sufficient power to detect hypothesized underlying latent classes.

Keywords: Bayesian information criterion, bootstrap likelihood ratio test, latent class analysis, mixture models, sample size, statistical power

In recent years, latent class analysis (LCA) has proven to be an important and widely used statistical tool in the social, behavioral, and health sciences. This technique, a form of finite mixture modeling (see McLachlan & Peel, 2000), can be used to identify underlying subgroups in a population. LCA can identify subgroups characterized by the intersection of particular behaviors, risk factors, or symptoms (e.g., Bucholz, Hesselbrock, Heath, Kramer, & Schuckit, 2000; Keel et al., 2004; Lanza, Rhoades, Greenberg, Cox, & The Family Life Project Key Investigators, 2011; Rindskopf & Rindskopf, 1986; Uebersax & Grove, 1990), such as symptoms of psychosis (e.g., Shevlin, Murphy, Dorahy, & Adamson, 2007), nicotine withdrawal symptoms (e.g., Xian et al., 2005), or adolescent risk behaviors (e.g., Collins & Lanza, 2010). Despite the benefits of the measurement model provided by LCA, several difficulties in applications of this method remain. One critical issue is that of model selection, especially selection of the number of classes, sometimes called *extraction* of classes (e.g., Nylund, Asparouhov, &

Muthén, 2007). The interpretation of all model parameters depends on the assumed number of classes. Simulations done to examine the performance of model selection tools in class extraction (e.g., Dziak, Coffman, Lanza, & Li, 2012; Nylund et al., 2007; Wu, 2009; Yang, 2006) have shown that having too small a sample often leads to choosing too few classes to adequately describe the data-generating model. However, to the best of our knowledge, there is no information in the literature about how to predict specifically how large a sample size N is actually needed to avoid such underextraction in practice.

In LCA, the failure to identify a substantively important, but perhaps not highly prevalent, latent class might lead to a violation of the local (conditional) independence assumption of LCA (i.e., the latent class variable is no longer adequate to fully describe the population covariance between the items; see Collins & Lanza, 2010). It might also cause a loss of important scientific information. For example, in the Shevlin et al. (2007) study, wrongly collapsing the hallucination (hallucinations only) and psychosis (multiple severe symptoms) classes together might give a misleadingly simplistic picture of the distribution of psychotic symptoms. Thus, statistical power for detecting latent classes can be as

Correspondence should be addressed to John J. Dziak, 204 E. Calder Way, Suite 400, State College, PA 16801, USA. E-mail: jjd264@psu.edu

important to the LCA user as statistical power for detecting significant effects is to the user of regression models. Although statistical power has been studied in the context of analysis of variance (ANOVA) and regression (e.g., Cohen, 1988) and in some covariance structure models (e.g., MacCallum, Browne, & Sugawara, 1996; MacCallum, Lee, & Browne, 2010; MacCallum, Widaman, Zhang, & Hong, 1999; Preacher & MacCallum, 2002; Satorra & Saris, 1985; Yuan & Hayashi, 2003), little is known about statistical power for detecting classes in LCA.

In this study we attempt to address this gap. First, we briefly review the bootstrap likelihood ratio test (BLRT), a very helpful procedure for testing hypotheses about the number of classes for LCA (see Nylund et al., 2007). Second, we briefly review how simulations can be used to construct power estimates for the BLRT given assumptions about the true population structure. Third, we propose effect size formulas based on the Cohen's w and Kullback–Leibler discrepancy measures. These formulas can be used in generalizing the results of our power simulations to new scenarios. Next, we provide extensive simulation results that show the usefulness of these effect size formulas. Finally, we provide tables and formulas for predicting required N for the BLRT in LCA and demonstrate their usefulness with additional simulations based on published latent class models. This work might help researchers decide how large a sample should be to have sufficient statistical power in tests for LCA class extraction. To our knowledge, power resources of this kind for the LCA BLRT were not previously available.

CHOOSING THE NUMBER OF CLASSES IN LCA

The LCA model for categorical observed items can be defined as follows. Let y_j represent element j of a response pattern \mathbf{y} . Let $I(y_j = r_j)$ equal 1 when the response to item j is r_j , and 0 otherwise. Then

$$P(\mathbf{Y}=\mathbf{y}) = \sum_{c=1}^K \gamma_c \prod_{j=1}^J \prod_{r_j=1}^{R_j} \rho_{j,r_j|c}^{I(y_j=r_j)}, \quad (1)$$

where γ_c is the probability of membership in latent class c and $\rho_{j,r_j|c}$ is the probability of response r_j to item j , conditional on membership in latent class c (see Lanza, Collins, Lemmon, & Schafer, 2007). The γ parameters represent the latent class membership probabilities. The ρ parameters represent item response probabilities conditional on latent class membership. The γ and ρ parameters can be estimated by maximum likelihood using the EM algorithm (Dempster, Laird, & Rubin, 1977).

The parameters of the model in Equation 1 cannot be estimated or interpreted without specifying a value of K . Sometimes theoretical reasons for choosing a particular K

are available, but more often researchers wish to use the data to guide their choice. They wish to avoid both underextraction (choosing a K that is too small) and overextraction (choosing a K that is too large). One approach is to compare models with 1, 2, 3, . . . latent classes, comparing the fit of each model to that of its predecessor using either a significance test or some relative fit criterion. If a K -class model fits significantly better than a $(K - 1)$ -class model, then the true population is assumed to have at least K classes.

The $(K - 1)$ - and K -class models can be compared by their best fitted log likelihoods. A naive approach involves comparing the likelihood ratio test (LRT) statistic $2l_{H1} - 2l_{H0}$, where l denotes the fitted log likelihood, to a χ^2 distribution with degrees of freedom $df = d_{H1} - d_{H0}$, where d denotes the number of parameters in the model. However, p values from this test are not valid because the usual asymptotic theory for the likelihood ratio statistic does not apply for comparing mixture models with different numbers of components. The parameter spaces of these models are not nested in the usual way,¹ and, as a result, p values from this naive test are not accurate (see Lin & Dayton, 1997; Lindsay, 1995; McLachlan & Peel, 2000). A more valid test can be obtained by the parametric BLRT, as further described by Feng and McCulloch (1996) and McLachlan and Peel (2000) and advocated by Nylund et al. (2007). The BLRT uses the same test statistic as the naive test, but it does not make the problematic assumption that the statistic has a χ^2 distribution under the null hypothesis. Instead, it uses an empirical estimate of the null hypothesis distribution.

The LCA BLRT can be performed as follows. First, fit the null $((K-1)$ -class) and alternative (K -class) models to the observed data set, and calculate the test statistic for each. Model identification must be examined for each latent class model being compared. In practice this requires multiple starting values. Next, using the parameter estimates from the null model, generate B separate random data sets. Simulation evidence in McLachlan (1987) suggests that B should be at least 99 to obtain optimal power; we use $B = 100$ in this article.² Now fit the null and alternative models to each generated data set and calculate the log likelihood for each null

¹We think of the $(K - 1)$ - and K -class models as conceptually nested, in that the former can be expressed as a special case of the latter (if any two classes have all ρ parameters equal or if any γ is zero). However, because of the lack of a unique representation of H_0 in the H_1 space, and because the H_0 values are on the boundaries of the space (e.g., a zero γ), classic asymptotic results for nested models do not hold (Lin & Dayton, 1997; Lindsay, 1995; McLachlan & Peel, 2000).

²We calculated the p value as the proportion of bootstrap data sets having a likelihood ratio test statistic greater than the corresponding statistic for the observed data set. Because we used 100 bootstrap data sets, we appropriately reject H_0 if $p \leq .05$ for $\alpha = .05$. If we had used 99 bootstrap data sets instead, which was recommended instead by Boos (2003), then it would have been appropriate instead to reject H_0 if $p < .05$ rather than $p \leq .05$. The distinction is important for maintaining the proper nominal Type I error rate, because bootstrap p values are fractions rather than continuous quantities.

and alternative model. Calculate the test statistic $2l_{H1} - 2l_{H0}$. To improve the chances of finding the global maximum likelihood solution for each model, multiple starting values should be used for each model here also.³ The empirical distribution of the B test statistics derived from the generated data sets can now serve as a reference distribution from which to calculate a critical value or a p value for the test statistic obtained from the observed data set.

In particular, let b be the number of generated data sets having calculated test statistics larger than the observed test statistic for the real data set. Then the bootstrap p value is calculated as b/B or $(b + 1)/(B + 1)$ (see Boos, 2003). The intuition is that if H_0 is true, then the observed data set will closely resemble the artificial data sets simulated under H_0 ; otherwise, it will be significantly different. Unlike the naive test, this test generally is approximately valid; that is, it empirically provides a true Type I error rate of approximately the specified nominal α or below when the null hypothesis is correct. Depending on the situation and the implementation of the test, it is possible for even a bootstrap test to have true Type I error rates slightly higher than nominal in some situations (e.g., see simulation results in McLachlan & Peel, 2000, p. 199) but it still does much better than a naive test would do (see Nylund et al., 2007, p. 554). There have been two main barriers to greater use of the BLRT. First, until recently, few software packages have provided it. However, it is implemented in *Mplus* (Muthén & Muthén, 2007), and a SAS macro for performing the BLRT in simple models using PROC LCA (Dziak, Lanza, & Xu, 2011; Lanza et al., 2007) is also available at <http://methodology.psu.edu/downloads/sasbootstrap>. Second, the BLRT can be computationally intensive. The LCA model must be fit multiple times (for different starting values) to each of B separate data sets, which could take minutes or even hours. Nonetheless, the BLRT is of special interest as the only widely used test for class extraction to provide a valid, known α level, so in this article we focus on the power of the BLRT.

The valid α level suggests that the risk of erroneously extracting too many classes can be controlled; that is, if the true number of classes is K and the model in Equation 1

holds, then size K will only be rejected 5% of the time or less in an $\alpha = .05$ test of size K versus size $K + 1$. However, even when a test is valid (i.e., controls Type I error), it could still have low power (fail to control Type II error). In the context of LCA, this means that although the BLRT is usually not very likely to extract too many classes (as its Type I error probability for a given comparison is very close to, or below, 5% for an $\alpha = .05$ test; see Nylund et al., 2007, p. 554) it is still likely that it might extract too few classes if the sample size is insufficient (see Nylund et al., 2007, p. 555). In the context of a test, underextraction of classes can be viewed as a lack of statistical power.

Consider the comparison of two models: a $(K - 1)$ -class model and a K -class model. By analogy to classic techniques, one could call the $(K - 1)$ -class model the null hypothesis and the K -class model the alternative hypothesis:

- H_0 : The population consists of $K - 1$ latent classes.
 H_1 : The population consists of K latent classes. (2)

The latent classes are assumed to be defined as in Equation 1. Thus we are considering Type I error here to denote choosing the K -class model when the $(K - 1)$ -class model is true, and Type II error to mean choosing the $(K - 1)$ -class model when the K -class model is true. Also by analogy to classic tests, the α level can be operationally defined as the probability of choosing the K -class model given a true $(K - 1)$ -class model, and “power” as the probability of choosing the K -class model given a true K -class model (i.e., one minus the Type II error probability). In an exploratory context, a comparison of a smaller to a larger model is likely to be done several times in sequence (1 vs. 2, 2 vs. 3, etc.) but for convenience we focus on a single step in this process.

AN EXAMPLE OF SIMULATED POWER FOR THE BLRT

As in classical settings (Cohen, 1988), power depends on N , as well as on the assumed parameters under both H_0 and H_1 . However, no theoretical formula, such as Cohen (1988) had for predicting the power of simpler tests, is available for predicting the power of the BLRT in LCA. Therefore, determining power here requires simulation. Consider the following example. Suppose that five forms of drug use (e.g., perhaps alcohol, tobacco, inhalants, cocaine, and prescription drug abuse) are being considered in a study of a particular population of at-risk youth, and participants are to be asked whether or not they had used each substance during the past year. These five questions correspond to five dichotomous items in an LCA. Suppose furthermore that in truth there are three classes. Members of the first class (“low use,” 60% of the population) have an independent 10% chance of endorsing each item. The second (“selective use,” 30%) has

³The minimum number of random starts needed in each sample or each bootstrap sample is not known. Using many starting values increases computational time, but using too few starting values, especially for the H_1 fits during the bootstrapping, can lead to invalid results due to suboptimal H_1 fits. In fact, if the H_1 estimates are based on a local maximum of the likelihood that is far from the global maximum, then the calculated likelihood ratio might occasionally even have a nonsensical negative value, indicating that the H_0 model space contains a higher likelihood peak than the larger H_1 space that includes it. If this problem were to occur frequently in simulations, then the power might be poorly estimated. With this in mind, we chose to use 50 random starts for H_0 and for H_1 in each of our simulations, both for the original simulated data sets and for each of their bootstrapped data sets. For the pseudo-population LCA used in calculating effect size, we used 200 random starts because this analysis seemed to be one in which precision was especially important, and because it only had to be done once per model and not separately for every simulated data set.

a 90% chance of endorsing the first three items, but only 10% of endorsing the last two. The third class (“high use,” 10%) has a 90% chance for each item. A researcher gathers a sample of size N from this population, without knowing the true class structure, and tests the (false) null hypothesis that there are two classes in the population versus the (true) alternative hypothesis that there are three. What would the power be for this test given N , or what N is required to obtain adequate power to select the true three-class model?

Power curves can be estimated by simulating data for a range of sizes N and comparing the results. For each sample size in $N = 50, 100, 150$, we simulated 1,000 data sets. In each, we fit the two-class and three-class models and compared them using a bootstrap test with $B = 100$ and $\alpha = .05$. For comparison, we also show the proportion selecting the three-class model as better than the two-class model if each of four other approaches are used: BLRT at a different α level, Akaike’s Information Criterion (AIC; Akaike, 1973), Schwarz’s Bayesian Information Criterion (BIC; Schwarz, 1978), or adjusted BIC (see Sclove, 1987). The proportion selecting the three-class model under each condition is shown in Table 1, which suggests that for the BLRT at a standard $\alpha = .05$, a power of about 80% would be achieved when N is slightly over 100. As expected, a higher or lower α is also associated with higher or lower power for the BLRT, respectively. The rejection probability for the two-class model using AIC is comparable to that of the $\alpha = .05$ bootstrap test in this case, and the rejection probability for the BIC is much lower. However, the rejection probabilities of the information criteria (IC) cannot really be interpreted as “power” here, because the IC are not intended as null hypothesis tests.

Of course, it is wise to consider multiple possible scenarios when planning a study. Thus, suppose we still assume a true H_1 with $K = 3$, but we are not sure about the class probabilities or the class-specific item response probabilities. Thus, we consider that the class memberships might be either approximately equal (34%, 33%, 33%) or unequal (60%, 30%, 10%). We also allow the class-specific responses to be .90 for high and .10 for low (a condition that we think of as strong measurement, or strong class separation), .80 and .20 (medium measurement), or .70 and .30 (weak

measurement). This leads to six scenarios arranged in a factorial design as shown in Table 2. The simulated power for each, with $N = 100$, shown in Table 3, suggests that measurement strength and relative class size both have drastic effects on power (Dziak, Min, & Lanza, 2009). Unequal class sizes and poor measurement led to a power of only 4% for the $\alpha = .05$ BLRT at $N = 100$, whereas the same test had a power of 99% under strong measurement and equal class sizes.

EFFECT SIZE FORMULAS

The findings in Table 3, as well as prior simulation work (e.g. Collins, Fidler, Wugalter, & Long, 1993; Lin & Dayton, 1997; Nylund et al., 2007; Wu, 2009), show that sample size requirements in LCA depend on several factors, such as relative class size and measurement strength. However, it is not clear how to use these ideas to generate a recommended minimum N for a particular situation. In general, estimation of power requires some assumptions about the true characteristics of the population or process being studied. For some simple tests, this is classically expressed in terms of an effect size measure (as in Cohen, 1988). An effect size measure can be useful both for prospective power analysis and perhaps also for quantifying the practical significance of obtained results. This simple idea might generalize to more complicated models.

For example, in the context of structural equation modeling (SEM) with multivariate normal data, it is known that for a given null hypothesis about the population covariance matrix, there can often be infinitely many sets of parameter values for H_1 offering exactly the same power against a given H_0 (MacCallum et al., 2010). This is a result of the higher dimensionality of the space in which the alternative hypothesis is defined (see MacCallum et al., 2010). However, although the situation appears more complicated than the hypotheses about single parameters considered by Cohen (1988), one can nevertheless still estimate power using a discrepancy measure based on the population difference in likelihood functions for the two population models. Such a population discrepancy measure is analogous to an effect size for a classic test, in that it can be used to solve for the power of the likelihood ratio test once df , N , and α are

TABLE 1
Simulated Power for a Three-Class Hypothetical Model for Different Sample Sizes

<i>N</i>	<i>Proportion Correctly Selecting 3-Class Over 2-Class Model</i>					
	<i>Using BLRT</i>			<i>Using Information Criteria</i>		
	$\alpha = .01$	$\alpha = .05$	$\alpha = .10$	<i>AIC</i>	<i>BIC</i>	<i>Adj. BIC</i>
50	0.24	0.44	0.55	0.37	0.04	0.94
100	0.49	0.72	0.81	0.77	0.11	0.91
150	0.73	0.90	0.94	0.93	0.21	0.95

Note. BLRT = bootstrap likelihood ratio test; AIC = Akaike information criterion (Akaike, 1973); BIC = Bayesian information criterion (Schwarz, 1978); Adj. BIC = adjusted BIC (see Sclove, 1987).

TABLE 2
Four Hypothetical Three-Class Models

Class label	Model 1			Model 2		
	No/Low	Med	High	No/Low	Med	High
Prob. of membership	.60	.30	.10	.34	.33	.33
Prob. of yes response						
Item 1	.30	.70	.70	.30	.70	.70
Item 2	.30	.70	.70	.30	.70	.70
Item 3	.30	.70	.70	.30	.70	.70
Item 4	.30	.30	.70	.30	.30	.70
Item 5	.30	.30	.70	.30	.30	.70

Class label	Model 3			Model 4		
	No/Low	Med	High	No/Low	Med	High
Prob. of membership	.60	.30	.10	.34	.33	.33
Prob. of yes response						
Item 1	.20	.80	.80	.20	.80	.80
Item 2	.20	.80	.80	.20	.80	.80
Item 3	.20	.80	.80	.20	.80	.80
Item 4	.20	.20	.80	.20	.20	.80
Item 5	.20	.20	.80	.20	.20	.80

Class label	Model 5			Model 6		
	No/Low	Med	High	No/Low	Med	High
Prob. of membership	.60	.30	.10	.34	.33	.33
Prob. of yes response						
Item 1	.10	.90	.90	.10	.90	.90
Item 2	.10	.90	.90	.10	.90	.90
Item 3	.10	.90	.90	.10	.90	.90
Item 4	.10	.10	.90	.10	.10	.90
Item 5	.10	.10	.90	.10	.10	.90

TABLE 3
Statistical Power for Several Hypothetical Models Having Three Latent Classes, Five Items, and $N = 100$

Model	Measurement Strength	Class Sizes	Proportion Correctly Selecting 3- Over 2-Class Model					
			BLRT			Information Criteria		
			$\alpha = .01$	$\alpha = .05$	$\alpha = .10$	AIC	BIC	Adj. BIC
1	Low	Unequal	0.01	0.04	0.08	0.08	0.00	0.29
2		Equal	0.01	0.05	0.11	0.10	0.00	0.35
3	Medium	Unequal	0.03	0.15	0.24	0.25	0.00	0.55
4		Equal	0.17	0.40	0.53	0.57	0.02	0.79
5	High	Unequal	0.49	0.72	0.81	0.77	0.11	0.91
6		Equal	0.94	0.99	1.00	1.00	0.73	1.00

Note. The model number in this table corresponds to the model number in Table 2. BLRT = bootstrap likelihood ratio test; AIC = Akaike information criterion (Akaike, 1973); BIC = Bayesian information criterion (Schwarz, 1978); Adj. BIC = adjusted BIC (see Sclove, 1987). Low, medium, and high “measurement strength” mean that ρ values are either .70 and .30, .80 and .20, or .90 and .10 for items typically endorsed or not endorsed, respectively, by members of each class, as presented in Table 2. Unequal and equal class sizes mean that the population class proportions are either different (10%, 30%, 60%) or approximately the same (34%, 33%, 33%) for members of each class.

specified (MacCallum et al., 1996; MacCallum et al., 2010; Satorra & Saris, 1985). The term *effect size* as we use it here is not intended to be interpreted as meaning the “effect” of an exogenous variable; we simply mean that it is a single

number that can act as a working definition of how far the null hypothesis is from the alternative hypothesis, for the sake of describing results or predicting power. We now seek such an effect size for the LCA BLRT as well, despite the

lack of multivariate normality and the presence of a mixture structure in LCA.

Cohen's w

We first propose that w , an effect size measure that Cohen (1988) proposed for classic Pearson χ^2 tests of hypotheses about contingency tables with m cells, could be used in a modified way for LCA as well. For a contingency table, let

$$w = \sqrt{\sum_{i=1}^{n_{\text{cells}}} \frac{(P_{1i} - P_{0i})^2}{P_{0i}}}, \quad (3)$$

where P_{0i} represents the probability of the i th cell, calculated under H_0 , and P_{1i} represents the probability of the i th cell, calculated under H_1 . Then w^2 has a form like a Pearson χ^2 test statistic, except that it uses assumed population proportions instead of sample proportions, and it does not directly depend on N . For an LCA model, the set of possible response profiles can be viewed as cells in a contingency table. For example, if there are six dichotomous items y_1, \dots, y_6 , then there are 2^6 possible response profiles $y = y_1, \dots, y_6$, and the probability of each of these profiles can be calculated using Equation 1 for H_0 and H_1 separately. Thus, w can be calculated given the hypothetical γ and ρ parameter values for the models corresponding to H_0 and H_1 .

For a classic χ^2 test of fit or of independence, without latent variables, Cohen's w is directly a function of the χ^2 noncentrality parameter and thus of power for a given df . Thus, in that specific context, if one knows N , df , the α level of the test, and the value of w , and if the asymptotic assumptions for the χ^2 test are met, power can be computed directly with no additional information required. For LCA applications, in which classic χ^2 tests are not necessarily applicable, some information could be lost by condensing all of the many γ and ρ parameters for H_1 and H_0 into the single value w , but as we later show, w is a fairly good predictor of approximate power.

The Kullback–Leibler Discrepancy

Similarly, we also propose the Kullback–Leibler (KL ; Kullback & Leibler, 1951) discrepancy between H_0 and H_1 as another useful effect size measure for the LCA BLRT. KL is a nonparametric measure of how different two fitted probability models are from each other. Denoting the models as M_a and M_b , the KL discrepancy can be written as

$$KL(M_a, M_b) = E_{M_a} \left(\log \left(\frac{L_{M_a}(y)}{L_{M_b}(y)} \right) \right), \quad (4)$$

where L is the fitted likelihood and E_{M_a} denotes the expectation computed under M_a . For many kinds of models, KL can be used to express how well a likelihood ratio test will

be able to detect the difference between the models, as it is the expected value of the log likelihood ratio (see Ferguson, 1996; van der Vaart, 1998). Using H_1 for M_a and H_0 for M_b , Equation 4 becomes

$$KL(H_1, H_0) = \sum_{i=1}^{n_{\text{cells}}} P_{1i} \left(\log \left(\frac{P_{1i}}{P_{0i}} \right) \right). \quad (5)$$

In calculating w or KL , we add a small constant such as 10^{-30} to each P_0 and P_1 to avoid dividing by 0 or taking a logarithm of 0 in the case of sparse data, although if the data are very sparse there could still be some numerical instability. Equation 5 cannot be re-expressed as a function of Equation 3, or vice versa; therefore, power predictions based on the two different effect size measures cannot be expected to match exactly. However, each measure is based on the degree of difference between the P_0 values and the P_1 values, and a value of zero for one effect size measure implies a value of zero for the other. w and KL are related to two classical ways of measuring the discrepancy between two probability models (the Pearson and Kullback–Leibler discrepancies, respectively, are further described in Zucchini, 2000), and they also closely resemble the Pearson χ^2 and likelihood or deviance G^2 LCA fit statistics (see Collins et al., 1993; Dayton, 1998), respectively. A brief outline of a convenient approach to calculating the w and KL effect size measures, requiring only a proposed set of γ and ρ parameters under H_1 , is given in Appendix A. We implement this approach in a SAS macro available at <http://methodology.psu.edu/LcaEffectSizes>.

Two Simpler Effect Size Formulas

Both the w and the KL measures already described include an expected cell proportion P_0 in the denominator for each possible cell. When there are J items, there are at least 2^J cells, so some expected cell proportions, and hence some denominators, will be extremely small. This raises questions about numerical stability and accuracy. Thus we next consider a measure based on the “index of dissimilarity” goodness-of-fit measure mentioned by Dayton (1998, p. 20). Revised somewhat for comparing two proposed population models rather than one proposed population model and a sample, it would be

$$I_D(H_1, H_0) = \frac{1}{2} \sum_{i=1}^{n_{\text{cells}}} |P_{1i} - P_{0i}|. \quad (6)$$

All three criteria considered so far involve P_{0i} and P_{1i} , the fitted probabilities for each cell of the contingency table under each hypothesized model. This is reasonable because these probabilities are the values that LCA models are intended to predict, in the same way that factor analysis models seek to predict covariance matrix entries (see Collins et al., 1993).

However, it is inconvenient in that there are many such probabilities (up to 2^J if each item is dichotomous, and more otherwise), some of which might be at or near zero. In a different context, sparse contingency tables are known to pose difficulty in interpreting χ^2 and G^2 fit statistics, which similarly involve a sum of ratios of tiny quantities over many cells (Collins et al., 1993; Muthén, 2008; Wu, 2009). Thus, one might wish for an effect size measure that would not directly involve the P_{0i} s and P_{1i} s.

From Wu (2009) and from Table 3, we know that having one or more extremely small classes (low γ values) or having one or more pairs of classes with similar response probabilities (similar ρ values) reduces power. Thus, among different possible parameter sets for a given model size (i.e., within a given combination of the number of items J and the number of H_1 classes K), $\min_{k=1,\dots,K} (\gamma_k)$ and $\min_{k,k'=1,\dots,K} \left(\sum_{j=1}^J (\rho_k - \rho_{k'})^2 \right)$ should each be positively related to power. Presumably, if one or more of the two similar classes was also a rare class, then there might be an even stronger tendency to lump the rare class into its more common neighbor. Thus, perhaps an ad-hoc measure of class separability such as

$$SEP = \min_{k,k'=1,\dots,K} \left(\gamma_k \gamma_{k'} \left(\sum_{j=1}^J (\rho_k - \rho_{k'})^2 \right) \right)$$

might be well correlated with power. This would be easier to calculate because it does not require P_0 or P_1 values. However, as we show later in our simulations, neither I_D nor SEP works as well for predicting power as w or KL .

SIMULATION EXPERIMENT 1: PERFORMANCE OF THE EFFECT SIZE FORMULAS

For the proposed effect size measures to be useful, they must first be shown to be predictive of power. Therefore, we did a simulation experiment as follows. For each of the $3 \times 2 = 6$ combinations of $J = 5, 9$, or 13 dichotomous items and of $K = 3$ or 5 classes, we created 10 random sets of γ and ρ parameters, each defining a possible joint population distribution of J dichotomous LCA items with no covariates. The probability ρ of a yes response on each item was generated as an independent uniform random variable for each class. Random γ parameters were generated such that for $k = 1, \dots, K$, $\gamma_k = e_k / \left(\sum_{j=1}^K e_j \right)$ where the e_j are independent from an exponential distribution with mean 1. For the simple case of $K = 2$, this method would let γ_1 be uniformly distributed between 0 and 1 (Casella & Berger, 1990; Jambunathan, 1954), and so γ_2 would be uniformly distributed as well because $\gamma_2 = 1 - \gamma_1$. Also, random ρ

parameters for the yes response were randomly drawn from a uniform distribution from 0 to 1. After this, to prevent extreme scenarios, parameter draws having any γ or ρ values less than .05 or ρ values greater than .95 were replaced with new random draws.

Within each of these $3 \times 2 \times 10$ random sets of parameters, we generated 100 data sets with $N = 500$ each and 100 data sets with $N = 1,000$ each, with the exception that in conditions with $J = 13$ and $N = 1,000$ only 50 data sets per parameter set were generated to save computation time. We also generated an extra 10 random sets of γ and ρ parameters for the $J = 5, K = 3$ condition, the $J = 9, K = 3$ condition, and the $J = 9, K = 5$ condition to obtain additional information; for each of these additional parameter sets we generated 50 data sets with $N = 500$ and 50 data sets with $N = 1,000$ per parameter set. These three conditions were of special interest because they were neither near-saturated like the $J = 5, K = 5$ condition, nor computationally burdensome like the $J = 13$ conditions. The total number of parameter sets was 200, and counting all of the 50 or 100 data sets per parameter set, the total number of data sets was 15,000. For each of these data sets, we performed a BLRT to compare the (false) H_0 that there were $K - 1$ classes to the (true) H_1 that there were K classes. Because each BLRT involved the original data set plus 100 bootstrap samples, and for each bootstrap sample the null and alternative model were fitted with 50 random starts each, the process required $15,000 \times (100 + 1) \times (50 + 50)$, or over 150 million LCA models to be fitted. The total computational cost was about 278 days distributed over several Linux processors.

The upper four panels of Figure A.1, in the online appendix, show scatterplots of the relationships of the four effect size measures considered (w , KL , I_D and SEP) with power in the $J = 5, K = 3$ case. Both w and KL seem to be strongly related to power. The other two measures are much poorer predictors of the power of the BLRT and are therefore not considered further. Figure A.1 also shows scatterplots of power against the Ramaswamy rescaled entropy statistic (Muthén, 2004; Ramaswamy, DeSarbo, Reibstein, & Robinson, 1993) for H_0 and H_1 .⁴ The entropy statistic measures how clearly or confidently a model classifies the

⁴To apply the idea of Ramaswamy entropy to a population model, we used the formula $1 + \frac{1}{\log K} \sum_{c=1}^K \sum_{i=1}^{2^J} P(\mathbf{y}_i) P(u_c | \mathbf{y}_i) \log P(u_c | \mathbf{y}_i)$, where \mathbf{y}_i is any of the 2^J possible response vectors, $P(\mathbf{y}_i) = \sum_{c=1}^K P(u_c) P(\mathbf{y}_i | u_c) = \sum_{c=1}^K \gamma_c P(\mathbf{y}_i | u_c)$ is a weighted average of the probability of \mathbf{y}_i across all classes calculated using the true parameters, and $P(u_c | \mathbf{y}_i)$ is the posterior probability of being in class c for an individual with data \mathbf{y}_i . (u_c here is an indicator for belonging to class c .) The usual Ramaswamy entropy formula in a sample is $1 + \frac{1}{N \log K} \sum_{c=1}^K \sum_{i=1}^N \hat{P}(u_c | \mathbf{y}_i) \log \hat{P}(u_c | \mathbf{y}_i)$ where the $\hat{P}(u_c | \mathbf{y}_i)$ are the posterior probabilities using the estimated parameters and where the \mathbf{y}_i are observed values from each subject. This is a reverse-signed and rescaled measure of entropy in the classical sense (Shannon, 1948); that is, of $-\sum \hat{P}(u_c | \mathbf{y}_i) \log \hat{P}(u_c | \mathbf{y}_i)$, so that higher Ramaswamy entropy means less unpredictability.

subjects in terms of posterior probabilities, but from Figure A.1 it is clear that neither H_0 entropy nor H_1 entropy by itself is a good effect size measure for a test; these plots appear to be mostly noise, with little or no systematic relationship with power.

Within each combination of J and K , the w and KL measures were strongly but not perfectly correlated ($r > .95$), except for the $J = 13, K = 3$ condition in which $r = .78$ due to an outlier. The outlier might have been due to a poorly calculated w caused by a near-zero denominator for one of the responses in one parameter draw. The range of values of w and KL that were observed differed by J and K . Because of the independent random ρ values, each item provided at least some evidence against H_0 (i.e., there were no items with exactly the same ρ value for two or more classes), so higher J was associated with higher w and KL and thus higher power. Also, roughly speaking, the larger K was, the more closely a $(K - 1)$ -class model could approximate a K -class model, so the effect size was lower.

It would be possible to interpolate the points for $N = 500$ or for $N = 1,000$ to get a point estimate of power for any value of the effect size measure (within the observed range) assuming the given value of N . However, it would be impractical to have to simulate separate power curves for each possible value of N . It would be more useful if the $N = 500$ and $N = 1,000$ curves could be combined with a single x axis and a single fit curve, allowing for interpolation and some extrapolation to new values of N . It seems plausible to assume additivity (e.g., that 500 independent observations from an H_1 with a discrepancy of $KL = 0.02$ from the H_0 would provide the same amount of evidence against H_0 as would 1,000 independent observations from an H_1 with $KL = 0.01$ from the H_0). Thus, perhaps $N \times KL$ can replace KL on the x axis of the plots and separate curves are not needed. The same approach might work with w , although it is necessary to square w before multiplying by N , because of the square root over the sum in Equation 3. Figures 1 and 2 show that additivity seems to apply, just as it would in a classic LRT.

An appealing feature of Figures 1 and 2 is that it is possible to characterize what kind of parameter set for H_1 is likely to be associated with a BLRT power of a given value (e.g., .80) at a given N for testing H_0 against H_1 . In the $J = 5, K = 3$ scenario, for example, it appears from Figure 1 that power of .80 would be obtained by parameter sets for which $N \times w^2 > 18$. Denoting this minimum $N \times w^2$ value as $m_{80}^{(w^2)}$, the researcher could choose

$$N = m_{80}^{(w^2)} / (w^2) \quad (7)$$

as being likely to detect a given true discrepancy of w . Equivalently, a true discrepancy of at least about $w = \sqrt{m_{80}^{(w^2)} / N}$ is needed to have 80% power to reject H_0 with




FIGURE 1 Scatterplot of $N \times w^2$ and simulated power for each simulated combination of number of items J and number of classes K . Empty circles represent $N = 500$ scenarios and solid circles represent $N = 1,000$ scenarios. The fit curves are obtained from a smoothing spline with roughness penalty coefficient chosen subjectively. A vertical dotted line marks the estimate of $m_{80}^{(w^2)}$, which we denote as the value of $N \times w^2$ for which power exceeds .80.

a sample size of $N = m_{80}^{(w^2)} = 18$ translates to a requirement for $N = 1,800$, $N = 200$, and $N = 72$, for $w = .1$, $.3$, and $.5$, respectively, the small, medium and large benchmark values used (in a different context) by Cohen (1988). A minimum $N \times w^2$ to obtain 90% power, although not specifically marked on the plot, could be obtained similarly and denoted $m_{90}^{(w^2)}$. Similarly, minimum $N \times KL$ benchmarks, denoted $m_{80}^{(KL)}$ and $m_{90}^{(KL)}$, can be estimated from Figure 2, and at least

$$N = m_{80}^{(KL)} / KL \quad (8)$$

is needed to obtain a power of at least .80 with a true discrepancy of KL of H_1 from H_0 . These estimates are shown in Table 4. However, w and KL were always near zero in the $J = 5, K = 5$ condition, because the H_1 models were almost saturated (29 parameters for a $2^5 = 32$ -cell




FIGURE 2 Scatterplot of $N \times KL$ and simulated power for each simulated combination of number of items J and number of classes K . Empty circles represent $N = 500$ scenarios and solid circles represent $N = 1,000$ scenarios. The fit curve is obtained in a similar way to that of Figure 1. A vertical dotted line marks the estimate of $m_{80}^{(KL)}$, which we denote as the value of $N \times KL$ for which power exceeds .80.

contingency table) and perhaps poorly identified, and so the values of $m_{80}^{(w^2)}$, and so on, could not be estimated for that condition.

It is clear from Table 4 that $m_{80}^{(w^2)}$ depends on J and possibly also on K . The dependence of power curves on model dimensionality is not surprising. The same holds for power in other covariance structure models (see, e.g., MacCallum et al., 2010; Saris & Satorra, 1993). Even in the simple case of a χ^2 test of fit or of independence on a contingency table (i.e., without latent variables, and with the assumption that N is adequate per cell for the appropriate χ^2 asymptotics to hold), Cohen (1988) still had to provide separate power tables depending on the df of the model. The underlying formula for the noncentrality parameter on which the power calculations were based was the same, but the reference noncentral χ^2 distribution depended on the df of the test.

For the LCA BLRT there is no reference χ^2 distribution. However, we can define an ad-hoc df for the test, as the number of distinct, freely estimated parameters in H_1 minus the number of distinct, freely estimated parameters in H_0 . The number of free parameters in Equation 1 with dichotomous items is $K - 1$ of γ parameters plus $J \times K$ of ρ parameters. Thus, for comparing an H_1 of K classes to an H_0 of $(K - 1)$ classes, the difference is $df = J + 1$; that is, 1 more γ parameter and J more ρ parameters. Recall that in a classic χ^2 test with high N and no latent variables, power can be computed using only N , w , and df (see Cohen, 1988). Therefore, we next conducted further simulations to test whether, in the LCA case, power could be adequately approximated using only $N \times w^2$ or $N \times KL$ and df (or equivalently J) and tabulated simulation results.

SIMULATION EXPERIMENT 2: CONSTRUCTING A TABLE OF BENCHMARK SAMPLE SIZE REQUIREMENTS

A thorough examination of the dependence of quantities such as $m_{80}^{(w^2)}$ on the number of items J and number of classes K requires additional simulations with more different values of J and K . Therefore, we performed new simulations on each combination of $J = 4, 5, 6, 7, 8, 9, 10, 11, 13, 15$ and $K = 2, 3, 4, 5, 6$. It was hypothesized that the shape of the power curves (i.e., the m values) would depend mainly on J and not K , although K might influence a point's position on the x axis (i.e., the effect size). If this conjecture is true, it would simplify the task of predicting power, because a useful reference table could be constructed by merely tabulating m values for every J in a range of interest, and not for every combination of J and K .

For each combination of J and K , we generated 10 random draws of parameters as described earlier. For each of these sets of random parameters, we simulated 200 data sets each with $N = 750$. For the easier scenarios ($J > 2K$), in which power might be high for every draw at $N = 750$, we additionally generated 200 more data sets having only $N = 250$ for each draw. To make this large number of simulations feasible, we employed a computational shortcut instead of doing the full BLRT each time. Resulting computational time was 554 days.⁵ Table 5 shows the median w and KL for

⁵For each J, K, N combination, we chose an estimated critical value as the average bootstrap critical value obtained from a preliminary study using only 10 bootstraps for each of 50 simulated data sets within each of 10 preliminary random draws of parameters. Because of the small number of bootstraps per data set, p values from these initial runs would be uninterpretable. However, we were only interested in obtaining an average value for the LRT statistic under H_0 for data sets and models of a certain size, not in a p value from each preliminary data set. We then obtained a power estimate from each of the sets of 200 data sets by doing tests similar to naive LRTs, except for using the estimated critical values obtained

TABLE 4
Estimated Constants m Based on Bootstrap Simulation Experiments

Number of Items	Number of Classes	Estimated $m_{80}^{(w^2)}$	Estimated $m_{90}^{(w^2)}$	Estimated $m_{80}^{(KL)}$	Estimated $m_{90}^{(KL)}$
Empirical estimates from Experiment 1					
5	3	17.9	21.4	8.5	10.3
9	3	34.0	40.8	13.7	16.6
9	5	35.1	43.3	13.5	16.5
13	3	62.3	79.0	18.1	21.6
13	5	57.1	70.9	16.7	20.7
Smoothed empirical estimates from Experiment 2					
4	Any	14.8	18.0	7.2	8.7
5	Any	16.2	19.9	7.1	8.6
6	Any	18.5	22.8	7.9	9.5
7	Any	22.0	27.1	9.1	11.0
8	Any	26.4	32.4	10.6	12.6
9	Any	30.8	38.0	11.7	14.1
10	Any	35.5	43.8	12.7	15.4
11	Any	40.4	49.6	13.6	16.3
12	Any	46.7	57.3	14.8	17.7
13	Any	54.6	67.3	16.2	19.2
14	Any	64.1	79.5	17.3	20.6
15	Any	74.3	92.7	18.3	21.7

Note. $m_{80}^{(w^2)}$ represents the constant in Equation 7 for predicting from w the required N to obtain a target power of .80, and $m_{90}^{(w^2)}$ represents the equivalent for a target power of .90. $m_{80}^{(KL)}$ represents the constant in Equation 8 for predicting from KL the required N to obtain a power of .80, and $m_{90}^{(KL)}$ represents the equivalent for a target power of .90. "Any" denotes a combined estimate for two through six classes.

TABLE 5
Median Effect Size Estimates for Testing K -Class Models Constructed Using Randomly Generated Parameter Draws, Versus $(K - 1)$ -Class Models, in Experiment 2

Number of Items	Median w Number of Classes					Median KL Number of Classes				
	2	3	4	5	6	2	3	4	5	6
4	0.153	0.050				0.011	0.001			
5	0.283	0.107	0.035			0.027	0.005	0.001		
6	0.438	0.114	0.057	0.039		0.057	0.006	0.002	0.001	
7	0.631	0.201	0.102	0.063	0.039	0.116	0.018	0.005	0.002	0.001
8	0.478	0.234	0.102	0.122	0.055	0.067	0.023	0.005	0.007	0.002
9	0.759	0.326	0.159	0.132	0.080	0.172	0.041	0.011	0.008	0.003
10	0.881	0.394	0.202	0.167	0.108	0.149	0.059	0.016	0.012	0.005
11	1.047	0.459	0.261	0.164	0.138	0.201	0.064	0.024	0.010	0.008
13	1.094	0.616	0.376	0.315	0.199	0.271	0.097	0.050	0.034	0.016
15	2.295	0.835	0.539	0.401	0.297	0.384	0.183	0.091	0.048	0.026

Note. Median rather than means are reported due to right skew.

the sets of parameters drawn under each combination of J and K . Figures 3 and 4 show how w and KL , respectively, are related to power within different values of J .

in this way in place of the naive χ^2 quantile. The goal was to construct a test hoped to have the same long-range power as the bootstrap test, although not necessarily the same result for a given data set. This shortcut was only used in Experiment 2, not Experiments 1 or 3. Because of the shortcut, the simulations for Simulation Experiment 2 required only 554 days of computational time despite being much more extensive than those of Simulation

Experiment 1. Most of this time was due to the larger models; specifically, if we had only implemented the conditions with $J \leq 6$ and $K \leq 4$, the total time would have been only 12 days.

Fit curves were applied to Figures 3 and 4 using nonparametric smoothing. Estimates for the m constants described earlier were obtained for each plot by simply recording where these curves crossed .80 or .90. These estimates, considered as a function of J , were then smoothed in an attempt




FIGURE 3 Scatterplot of $N \times w^2$ and simulated power for models based on different random parameter draws in Experiment 2. The plot symbols for each simulated model are numerals representing the number of classes for that model. The fit curves are obtained from a smoothing spline with roughness penalty coefficient chosen subjectively. A vertical dotted line marks the $m_{80}^{(w^2)}$ estimate.

to reduce random error and allow interpolation to the $J = 12$ and $J = 14$ cases, which had been omitted to save computer time.⁶ The resulting estimates are shown in the second part of Table 4.

⁶The nonparametric smoothing for the plots and for Table 4 was done using smoothing splines, implemented either with the R function `smooth.spline` (R Development Core Team, 2010) or with the R function `gam` in library `mgcv` (Wahba, 1990; Wood, 2003, 2011, 2012).

Especially for *KL*, a single curve for each J did fairly well in summarizing the relationship of effect size to power regardless of K . Although this does not prove that K does not affect the power curve, it does suggest that a table of $m_{80}^{(KL)}$ constants indexed only by J might be useful. The curve (i.e., the power for a given effect size) did not depend noticeably on K , although the position of points on the curve (i.e., the effect size actually obtained) did depend heavily on K . Table 5 shows that the effect sizes obtained depend jointly on J and K , with highest effect sizes observed when J is large and K is small.




FIGURE 4 Scatterplot of $N \times KL$ and simulated power for models based on different random parameter draws in Experiment 2. The plot symbols for each simulated model are numerals representing the number of classes for that model. The fit curves are obtained from a smoothing spline with roughness penalty coefficient chosen subjectively. A vertical dotted line marks the $m_{80}^{(KL)}$ estimate.

SIMULATION EXPERIMENT 3: VALIDATING ESTIMATED SAMPLE SIZE REQUIREMENTS USING EMPIRICAL MODELS

The plots from Experiment 2 provide sufficient information to make power predictions for a wide range of LCA scenarios. However, there are still some reasons to be uncertain about the performance of these predictions. First, they are based on a computational shortcut. The m estimates for a given J (e.g., $J = 9$) in Table 4 seem to be slightly lower on average for comparable conditions in Experiment 2 (with the shortcut) than in Experiment 1 (without the shortcut). Thus,

it is not clear whether the shortcut leads to underestimating the needed N . Second, both Experiments 1 and 2 were based on randomly generating parameter values, and therefore might not be well representative of the parameter values that would be found in real-world research.

To address these concerns about validity, we performed a final set of simulation experiments. For each of the six parameter sets in Table 2, and additionally for each of seven empirical examples of fitted LCA models chosen from recent literature, we first imagined that the proposed or estimated model was exactly true for the population. We imagined that an investigator was planning to gather

TABLE 6
Agreement With Target Power for K -Class Versus $(K - 1)$ -Class Models in Experiment 3

Scenarios	Items	Classes	Effect Sizes		Using N_w		Using N_{kl}		Using $1.15 N_{kl}$	
			w	KL	N	Power	N	Power	N	Power
Scenarios from hypothetical models										
Model 1	5	3	0.06	0.002	4071	0.82	3571	0.76	4107	0.82
Model 2	5	3	0.11	0.006	1314	0.77	1163	0.72	1338	0.78
Model 3	5	3	0.17	0.014	559	0.77	494	0.74	569	0.79
Model 4	5	3	0.27	0.038	216	0.81	188	0.74	217	0.80
Model 5	5	3	0.35	0.061	131	0.87	117	0.82	135	0.88
Model 6	5	3	0.51	0.141	62	0.93	51	0.87	59	0.91
Scenarios from published empirical results										
Alcohol	7	5	0.27	0.040	312	1.00	230	0.97	265	0.98
Delinquency	6	4	0.18	0.015	586	0.77	531	0.71	611	0.79
Depression	8	5	0.23	0.025	491	0.79	428	0.69	493	0.79
Eating	14	4	0.58	0.140	191	0.98	124	0.87	143	0.94
Health risks	12	5	0.30	0.028	537	0.77	534	0.77	615	0.85
Nicotine	12	4	0.18	0.014	1524	0.95	1075	0.78	1237	0.86
Psychosis	13	4	0.11	0.005	4639	0.98	3219	0.84	3702	0.87

Note. w and KL refer to the calculated effect sizes for each model. N_w , N_{KL} , and $1.15 N_{KL}$ refer to sample size N chosen using each of the three ways described in the text for calculating needed sample size (using w , using KL , or using KL and then multiplying the result by 1.15). For each combination of model and sample size calculation method, N refers to the sample size recommended by that method to obtain a predicted power of .80, and Power refers to the simulated power actually achieved for samples of that size from that model. Thus, a highly accurate method of predicting power would have entries very close to .80 in its power column.

data from this population to test whether there were K or $K - 1$ classes, where K was the number of classes in the published work. We used the following empirical examples. The *alcohol* model is taken from an analysis in Lanza et al. (2007, p. 687) of data from Monitoring the Future (Johnston, Bachman, O'Malley, & Schulenberg, 2004), original $N = 2,490$. The *delinquency* model is from Collins and Lanza (2010, p. 12), original $N = 2,087$, based on Wave I data from the National Longitudinal Study of Adolescent Health (Udry, 2003). The *depression* model is from Lanza, Flaherty, and Collins (2003, p. 687), original $N = 1,892$. The *eating* model is from Lanza, Savage, and Birch (2010, p. 836), original $N = 197$. The *health risks* model is from Collins and Lanza (2010, p. 39), original $N = 13,840$, using Youth Risk Behavior Survey (2005) data. The *nicotine withdrawal* model is from Xian and colleagues (2005, p. 414), original $N = 4,112$. The *psychosis* model is from Shevlin, Murphy, Dorahy, and Adamson (2007, p. 105), original $N = 5,893$.⁷ Each of these models is based on dichotomous items, typically representing the presence or absence of some symptom or behavior, as is fairly common in psychological and medical research.

For each data set, we calculated w and KL based on the population parameters, and then calculated the required N

according to Table 4 to obtain a power of .80 at $\alpha = .05$ for that w or KL (i.e., we imagined that the hypothetical experimenter guesses the effect size exactly right). We then simulated 1,000 data sets of size N from this hypothetical population, and performed the bootstrap test (without the computational shortcut) for each one. We did this separately for N_w , the sample size obtained using w and Equation 7, and for N_{KL} , the sample size obtained using KL and Equation 8. Finally, we also did a set of simulations for $1.15N_{KL}$, for reasons described later. The total computational time required was 221 days. Results are shown in Table 6.

Ideally, we would hope that the simulated power for each of these 13 scenarios would be close to .80. This would indicate that the results of Experiment 2 were generalizable to parameter sets other than those used in Experiment 2. Table 6 shows that the sample sizes chosen to afford a power of .80 do obtain power reasonably near .80, although how near seems to depend on the method as well as the scenario. N_w is seldom too small (providing power of at least .77 for all of the models) but often needlessly large (power higher than .99 in one instance). N_{KL} provides power that is very close to .80 on average, and seems to be a more precise estimate than the estimate from w . Thus, KL might potentially be a more useful measure than w for sample size planning.

However, there is an important caveat: Notice that N_{KL} is sometimes too small (in one case a power of only .69 is obtained for a target power of .80). In many cases, an underestimation of the needed sample size could be more harmful than an overestimation (i.e., it is better to err by expending slightly more resources than is needed to attain one's goals,

⁷The ρ estimates in this article were provided as a plot rather than a table, so we estimated them by eye from the plot. We wished to use this model partly because of its interesting composition: As an epidemiological survey of a severe condition, it had very unbalanced class sizes (only 2% were in the severest class).

TABLE 7
 H_0 Rejection Probability of Different Methods in Experiment 3, Using N Determined From w and Equation 7

		Proportion Correctly Selecting K -Class Over $(K-1)$ -Class Model					
		BLRT With			Information Criteria		
	N	$\alpha = .01$	$\alpha = .05$	$\alpha = .10$	AIC	BIC	Adj. BIC
Hypothetical scenarios							
Model 1	4,071	0.58	0.82	0.87	0.92	0.01	0.17
Model 2	1,314	0.50	0.77	0.86	0.91	0.02	0.40
Model 3	559	0.57	0.77	0.87	0.90	0.05	0.60
Model 4	216	0.55	0.81	0.89	0.92	0.13	0.88
Model 5	131	0.67	0.87	0.92	0.92	0.19	0.96
Model 6	62	0.78	0.93	0.96	0.94	0.38	1.00
Scenarios from published empirical results							
alcohol	312	0.89	1.00	1.00	0.89	0.00	0.67
delinquency	586	0.54	0.77	0.83	0.88	0.02	0.54
depression	491	0.57	0.79	0.87	0.95	0.00	0.62
eating	191	0.94	0.98	0.99	0.99	0.08	0.98
health risks	537	0.57	0.77	0.85	0.96	0.00	0.56
nicotine	1,524	0.84	0.95	0.97	1.00	0.00	0.49
psychosis	4,639	0.92	0.98	0.99	1.00	0.00	0.24

Note. BLRT = bootstrap likelihood ratio test; AIC = Akaike information criterion (Akaike, 1973); BIC = Bayesian information criterion (Schwarz, 1978); Adj. BIC = adjusted BIC (see Sclove, 1987).

than to fail to attain them at all by expending slightly too little). Thus, perhaps an unbiased estimate of the required sample size is not really desirable, but instead it would be better to have an estimate that is purposefully slightly biased in the direction of caution. This might be done by first obtaining an N estimate from KL but then multiplying this estimate by some constant slightly greater than 1. We found that the multiplier 1.15 seems to work well,⁸ in that the selected N s now provided simulated power ranging from about .8 to about .9. That is, using $1.15N_{KL}$ gave power values that were never much lower than .80, although they were often slightly higher.

The correct selection rates by other analysis methods besides the $\alpha = .05$ bootstrap were also observed, and are compared in Table 7. For simplicity we show only results using Nw because relative patterns will be the same regardless. The values for the BLRT in Table 7 can be considered to represent statistical power. The term *power* is not quite correct for AIC, BIC, or adjusted BIC because they are not designed to be hypothesis tests with fixed α levels (although they are being used like tests in this case, in that a larger model is being compared to a smaller). However, it is still enlightening to compare their behavior. Decreasing α for the BLRT from .05 to .01 noticeably reduces power, although increasing α to .10 does not increase it as much. The rejection rates of the too-small model by the IC varied,

with AIC being the most liberal or sensitive (hence least likely to result in underextraction) and BIC being by far the most conservative, parsimonious, or specific (most likely to result in underextraction). Although this is important information, decisions about what selection method to use should not be made on the basis of Table 7 alone, because it only deals with the case in which H_1 is true. If overextraction (the equivalent of Type I error) were being considered, then the reverse would have been seen, with BIC becoming more likely to be correct than AIC, and a BLRT with a low nominal α becoming more likely to be correct than one with a higher one (see Dziak et al., 2012).

RECOMMENDATIONS

Using the results in Table 4, a predicted sample size requirement can be obtained for powering a test to select the K -class model over the corresponding $(K - 1)$ -class model, at least under the assumption that there are no more than K classes. This could be useful in planning the needed N for observational studies based on consideration of one or more hypothetical population structures. It might also be useful in evaluating the adequacy of available data sets prior to doing an LCA. Using KL seems to be a slightly better choice than using w , but w also works fairly well. To err on the side of caution, we recommend multiplying the estimated required N by 1.15.

As illustrated in Table 8, recommended N depends heavily on the assumed effect size. Cohen (1988) provided benchmarks of .1, .3, or .5 for low, medium or high w , but these were for a classic χ^2 fit test, not for LCA. We are unaware of any similar benchmarks for KL at all. The average values in

⁸We originally chose the value 1.15 partly because we observed that most of the $m_{80}^{(KL)}$ values from the upper part of Table 4 (obtained without the computational shortcut) were between 1.1 and 1.2 times their counterparts in the lower part of the table (obtained with the shortcut), so that a multiplier of 1.15 might counteract any bias toward underestimation that might have been caused by the shortcut.

TABLE 8
Sample Size Recommendations for Power = .80 Assuming Different Values of w

<i>J</i>	<i>Assuming Benchmark Values of w</i>			<i>Assuming Median Observed w From Random Data Sets in Experiment 2 Given a Number of Classes</i>									
	<i>(w = .1)</i>	<i>(w = .3)</i>	<i>(w = .5)</i>	<i>(2 Classes)</i>		<i>(3 Classes)</i>		<i>(4 Classes)</i>		<i>(5 Classes)</i>		<i>(6 Classes)</i>	
	<i>N</i>	<i>N</i>	<i>N</i>	<i>w</i>	<i>N</i>	<i>w</i>	<i>N</i>	<i>w</i>	<i>N</i>	<i>w</i>	<i>N</i>	<i>w</i>	<i>N</i>
4	1,480	164	59	0.15	632	0.05	5,920	Low	High	Low	High	Low	High
5	1,620	180	65	0.28	202	0.11	1,415	0.04	13,224	Low	High	Low	High
6	1,850	206	74	0.44	96	0.11	1,424	0.06	5,694	0.04	12,163	Low	High
7	2,200	244	88	0.63	55	0.20	545	0.10	2,115	0.06	5,543	0.04	14,464
8	2,640	293	106	0.48	116	0.23	482	0.10	2,537	0.12	1,774	0.06	8,727
9	3,080	342	123	0.76	53	0.33	290	0.16	1,218	0.13	1,768	0.08	4,813
10	3,550	394	142	0.88	Low	0.39	229	0.20	870	0.17	1,273	0.11	3,044
11	4,040	449	162	1.05	Low	0.46	192	0.26	593	0.16	1,502	0.14	2,121
13	5,460	607	218	1.09	Low	0.62	144	0.38	386	0.32	550	0.20	1,379
15	7,430	826	297	2.30	Low	0.84	107	0.54	256	0.40	462	0.30	842

Note. J represents the number of items. K represents the number of classes for the recommendations that use median values from Experiment 2. The number of classes does not matter for the recommendations that simply use the benchmark value. The estimated required sample size N for a given value of w was calculated using Equation 7 and the $m_{80}^{(w^2)}$ constants from Table 4 here, without multiplying by our proposed conservative factor of 1.15. To apply the conservative factor, simply multiply the N values by 1.15.

Table 5 could be used, but these estimates might be too pessimistic because they are based on random numbers rather than interesting and well-measured real-world data sets.

Table 6 might provide a better guide to reasonable effect size values than Table 5, as it is based on published estimates rather than random data. However, in some cases it might be better for a researcher to construct one or more scenarios deemed reasonable for the true parameters, as in Table 2, and then calculate w or KL from each of them using the pseudo-population approach described earlier in the article; a macro for doing this is provided at <http://methodology.psu.edu/LcaEffectSizes/>.

In summary, the required sample size for an LCA BLRT can be estimated as follows. First, specify an estimated effect size (either w or KL). One way to obtain such an estimate is to specify an H_1 model (including the number of items, supposed number of classes, and supposed γ and ρ values), and calculate the effect size for testing this model against the H_0 number of classes (a SAS macro is available for easily doing so in the case of comparing a K vs. $K - 1$ class model). A conservative approach would be to consider several reasonable models and choose the smallest effect size among them. After obtaining the effect size, one can use Equation 7 or 8 to estimate the required N without needing to do simulations. If desired, one could optionally justify this estimate further by simulating data sets with that N and calculating the power as we did in Experiment 3.

DISCUSSION

In this article we empirically addressed the question of how large a sample size is needed to avoid underextraction when using the BLRT to choose a number of classes in LCA. We proposed a method for doing so, based on

first specifying a proposed effect size in terms of w or KL , and then using the tables compiled from our simulations to calculate the power needed to detect that effect size with adequate probability. Our simulations involved generating many scenarios with random characteristics, calculating the effect size measures, and empirically estimating the power of the BLRT on many simulated data sets for each scenario. This provided a feasible, although time-consuming, way to get around the lack of an analytic power formula for the LCA BLRT: namely, empirically plotting simulated power against simulated effect size to tabulate the required sample sizes. The tables can now be used by investigators without having to do simulations themselves. We studied only the case of dichotomous items, but the approach could be extended in the future to more general data sets.

An obvious limitation of this article is that we considered only a fairly simple application of LCA, involving dichotomous items and no covariates. Our study of even this simple case required a very large investment of computation time, so it is not clear how future research on more complicated models should be done. It is not known how polytomous items would affect power. It is reasonable to expect that informative covariates might improve power (Wu, 2009, p. 116), but this would require further study, as would the question of predicting power to detect relationships between covariates and the latent class variable.

Another important caveat is that the results in this article are better used for prospective planning prior to doing an analysis, and might not be meaningful if applied in a post-hoc way after doing an LCA (just as with classical power and sample size formulas; see, e.g., Hoenig & Heisey, 2001; Lenth, 2007). For example, in Table 6, the required N was always calculated as less than the N that was used in the original study. However, this cannot be interpreted as

evidence that the original studies successfully found “all” of the classes of interest in their population, even if the classes are considered to be real phenomena and not just convenient abstractions in a model. Similarly, it cannot be interpreted as evidence that past LCA studies have used large enough N , or too large N , according to some objective standard. Either conclusion would involve circular reasoning, because by doing the computation we had already assumed that the results of the original study were accurate. Thus, although our results might be useful in planning empirical studies, or even in deciding whether to do an LCA on an existing data set, they should not be used to attempt to quantify confidence in a fitted LCA model.

Also, although we have treated model selection as closely related to testing, they are not the same (see, e.g., Burnham & Anderson, 2002). By focusing on testing in this article, we gloss over some complications, especially the possibility that neither model might be correct. This is relevant even in classic significance testing, but is very obvious in the context of class extraction. Nonetheless, the familiar paradigm of the power of a statistical test provides a simple and tractable way to begin to explore required N . In practice, researchers might take a stepwise approach to LCA model selection, such as comparing a one-class solution to a two-class, then a two-class solution to a three-class, and so on. If a researcher compares each pair of model sizes with a penalized likelihood criterion such as AIC (Akaike, 1973) or BIC (Schwarz, 1978), choosing the best one according to this criterion, then this is algebraically exactly the same as comparing each pair of models using a naive likelihood ratio test, with an α level determined indirectly by the penalty weight of the criterion (see Foster & George, 1994; Leeb & Pötscher, 2005; Teräsvirta & Mellin, 1986). Thus, our approach based on ideas of hypothesis testing seems to be reasonable. However, there is not enough information to comprehensively compare approaches such as the IC to the BLRT. In Tables 1, 3, and 7, it is clear that AIC has a much higher probability of rejecting the inadequate model than BIC does. Reliance on the BIC is likely to result in the choice of overly simple models if N is not large.⁹ However, we do not necessarily recommend AIC as a substitute for BIC, because that might

replace the underextraction problem with an overextraction problem. It is known that when using AIC or BIC to choose between two models, AIC will be more likely to overfit and BIC will be more likely to underfit (Dziak et al., 2012; Lin & Dayton, 1997). Rather than choosing one extreme or the other, it might be better to consider both AIC and BIC together (Collins & Lanza, 2010). Using BIC alone, however, involves a serious underfitting risk.¹⁰

We have only explored tests of the form of Equation 2; that is, $K - 1$ classes versus K classes. We have not explored what would happen if we replace H_0 in Equation 2 with “The population consists of $K - 1$ or fewer latent classes.” However, there might not be much of a difference. That is, if a $(K - 2)$ -class model is not rejected by the BLRT against a K -class alternative, then a $(K - 1)$ -class model should not be rejected either, assuming enough random starts have been used to find a global optimum likelihood for each size. A more fundamentally different case is what to do if one wishes to test broader alternative hypotheses. For example,

H_0 : The population consists of at most $K - 1$ latent classes.

H_1 : The population consists of K or more latent classes. (9)

Such hypotheses might be of interest in confirmatory or goodness-of-fit analyses with LCA. However, the approach taken in this article does not apply well to them because of the lack of a specific H_1 . The test in Equation 9 is best viewed as a test of goodness of fit of the $K - 1$ class model; that is, of whether the local independence assumption implicit in Equation 1 holds (see Collins et al., 1993; Collins & Lanza,

advantage over BIC because they considered only the possibility of underextraction and not of overextraction, whereas the Nylund paper considered both (see their Table 7).

¹⁰Whether it is desirable to have high power when effect sizes are very small is a deeper question; that is, whether one should wish to reject models that have statistically significant lack of fit but are practically interpretable, in favor of models that quantitatively describe the population covariance structure slightly more accurately but are harder to interpret (see Dziak et al., 2012; Raftery, 1995). This would depend on the researcher’s goals. We have assumed that Equation 1 was literally true, and hence that the classes were truly distinct “real” groups, not just summaries of an underlying continuum; this was necessary to have an unambiguous true value of K in the simulations. However, it is often unclear whether latent constructs should be seen as categorical or continuous, and this might depend on the context (see, e.g., Bauer & Curran, 2003, with responses and rejoinder; Collins & Lanza, 2010; Walters, 2011). If the indicators are really a function of a continuous rather than categorical latent variable, Equation 1 will presumably be rejected for high enough N , at any $K < 2^J$, as its local independence assumption would not be precisely true at any reasonable number of classes. The idea of a correct K could perhaps be replaced by optimizing some quantitative measure of estimation quality, but the best K would grow with larger N due to a changing bias–variance trade-off, and would likely represent a less parsimonious and interpretable model than might otherwise be wished; AIC might be more appropriate than BIC or the bootstrap under this alternative viewpoint (see Burnham & Anderson, 2002; Dziak et al., 2012). However, if a small and easily interpreted model were desired, BLRT or BIC might point out a better model size than AIC.

⁹Nylund and colleagues (2007) concluded from their simulations that BIC usually performs better than AIC, especially when N is large. However, our results suggest that BIC sometimes performs extremely poorly with modest N because of a very high probability of extracting too few classes. A partial explanation of the difference between the findings is in the models considered: The models considered by Nylund and colleagues had very good measurement and class separation. In the context of our articles, they used rather large effect sizes: The effect sizes for the four categorical-outcome LCA examples in their Table 2 were $w = 0.80, 0.44, 3.97, 0.66$, respectively or $KL = 0.23, 0.10, 1.31, 0.13$, respectively, higher than most of the values in our Tables 5 or 6. In a sense, this gave BIC an unfair advantage in Nylund and colleagues (2007), because models were not considered in which the $K - 1$ class model was somewhat wrong but not grossly wrong (i.e., in which there was a true but subtle effect to which AIC would be more sensitive). On the other hand, our simulations admittedly gave AIC an unfair

2010; Vermunt & Magidson, 2004, 2005), rather than as a comparison of two specific candidate values of K . Thus, because the test is different, the appropriate procedure for power or sample size planning would presumably also have to be different, and might be based on trying to determine the power to detect a certain amount of residual correlation.

Despite these limitations, in this study we have found that two proposed effect size measures, w and KL , were strongly related to the power of the LCA BLRT across thousands of simulated data sets. Based on this, we have compiled power tables for LCA class extraction, which has not been done before to our knowledge. We hope that this will help applied researchers in planning studies when latent classes are posited to exist.

ACKNOWLEDGMENTS

We thank Dr. YoungKyoung Min for extensive assistance in implementing important Monte Carlo studies during an earlier phase of the project. We thank Drs. Beau Abar, Donna Coffman, and Mildred Maldonado-Molina for reviewing an early draft of this article, and Dr. Runze Li for suggestions about computational resources. We thank Amanda Applegate for her very helpful review and proofreading.

Effect size computations were done in SAS using PROC LCA. SAS software is copyright 2002–2012 by SAS Institute, Inc. SAS and all other SAS Institute, Inc. product or service names are registered trademarks or trademarks of SAS Institute Inc., Cary, NC, USA. The simulations themselves were done in R (R Development Core Team, 2010) using a special adaptation of the computational engine of PROC LCA (Lanza, Lemmon, et al., 2010) adapted by Xianming Tan to run in R on Linux.

FUNDING

This research was supported by NIDA grant DA010075-15. The content is solely the responsibility of the authors and does not necessarily represent the official views of the NIDA or the National Institutes of Health.

REFERENCES

- Akaike, H. (1973). Information theory and an extension of the maximum likelihood principle. In B. N. Petrov & F. Csaki (Eds.), *Second international symposium on information theory* (pp. 267–281). Budapest, Hungary: Akademiai Kiado.
- Bauer, D. J., & Curran, P. J. (2003). Distributional assumptions of growth mixture models: Implications for overextraction of latent trajectory classes. *Psychological Methods*, 8, 338–363.
- Boos, D. D. (2003). Introduction to the bootstrap world. *Statistical Science*, 18, 168–174.
- Buchholz, K., Hesselbrock, V., Heath, A., Kramer, J., & Schuckit, M. (2000). A latent class analysis of antisocial personality disorder symptom data from a multi-centre family study of alcoholism. *Addiction*, 95, 553–567.
- Burnham, K. P., & Anderson, D. R. (2002). *Model selection and multi-model inference: A practical information-theoretic approach* (2nd ed.). New York, NY: Springer-Verlag.
- Casella, G., & Berger, R. L. (1990). *Statistical inference*. Belmont, CA: Wadsworth.
- Cohen, J. (1988). *Statistical power analysis for the behavioral sciences* (2nd ed.). Hillsdale, NJ: Erlbaum.
- Collins, L. M., Fidler, P. L., Wugalter, S. E., & Long, J. D. (1993). Goodness-of-fit testing for latent class models. *Multivariate Behavioral Research*, 28, 375–389.
- Collins, L. M., & Lanza, S. T. (2010). *Latent class and latent transition analysis: With applications in the social, behavioral, and health sciences*. New York, NY: Wiley.
- Dayton, C. M. (1998). *Latent class scaling analysis* (Sage University Papers Series on Quantitative Applications in the Social Sciences, No. 07–126). Thousand Oaks, CA: Sage.
- Dempster, A. P., Laird, N. M., & Rubin, D. B. (1977). Maximum likelihood from incomplete data via the EM algorithm. *Journal of the Royal Statistical Society, Series B*, 39, 1–38.
- Dziak, J. J., Coffman, D. L., Lanza, S. T., & Li, R. (2012). *Sensitivity and specificity of information criteria* (Tech. Rep. No. 12–119). University Park, PA: The Pennsylvania State University, The Methodology Center. Retrieved from <http://methodology.psu.edu>
- Dziak, J. J., Lanza, S. T., & Xu, S. (2011). *LcaBootstrap SAS macro users' guide* (Version 1.1.0). University Park, PA: The Methodology Center, The Pennsylvania State University. Retrieved from <http://methodology.psu.edu>
- Dziak, J. J., Min, Y. K., & Lanza, S. T. (2009, May). *Sample size requirements in latent class analysis: The power of the bootstrap test for detecting a class*. Poster presented at the annual meeting of the Society for Prevention Research, Washington, DC.
- Feng, Z. D., & McCulloch, C. E. (1996). Using bootstrap likelihood ratios in finite mixture models. *Journal of the Royal Statistical Society, Series B*, 58, 609–617.
- Ferguson, T. S. (1996). *A course in large sample theory*. London, UK: Chapman & Hall.
- Foster, D. P., & George, E. I. (1994). The risk inflation criterion for multiple regression. *Annals of Statistics*, 22, 1947–1975.
- Hoenig, J. M., & Heisey, D. M. (2001). The abuse of power: The pervasive fallacy of power calculations in data analysis. *The American Statistician*, 55, 19–24.
- Jambunathan, M. V. (1954). Some properties of beta and gamma distributions. *The Annals of Mathematical Statistics*, 25, 401–405.
- Johnston, L. D., Bachman, J. G., O'Malley, P. M., & Schulenberg, J. E. (2004). *Monitoring the future: A continuing study of American youth (12th-grade survey)* [Machine-readable data file and documentation]. Ann Arbor, MI: Inter-University Consortium for Political and Social Research.
- Keel, P. K., Fichter, M., Quadflieg, N., Bulik, C. M., Baxter, M. G., Thornton, L., Halmi, K. A., Kaplan, A. S., Strober, M., Woodside, D. B., Crow, S. J., Mitchell, J. E., Rotondo, A., Mauri, M., Cassano, G., Treasure, J., Goldman, D., Berrettini, W. H., & Kaye, W. H. (2004). Application of a latent class analysis to empirically define eating disorder phenotypes. *Archives of General Psychiatry*, 61, 192–200.
- Kullback, S., & Leibler, R. A. (1951). On information and sufficiency. *Annals of Mathematical Statistics*, 22, 79–86.
- Lanza, S. T., Collins, L. M., Lemmon, D. R., & Schafer, J. L. (2007). PROC LCA: A SAS procedure for latent class analysis. *Structural Equation Modeling*, 14, 671–694.
- Lanza, S. T., Flaherty, B. P., & Collins, L. M. (2003). Latent class and latent transition analysis. In J. A. Schinka & W. F. Velicer (Eds.), *Handbook of psychology* (Vol. 2, pp. 663–665). Hoboken, NJ: Wiley.

- Lanza, S. T., Lemmon, D. R., Dziak, J. J., Huang, L., Schafer, J. L., & Collins, L. M. (2010). *PROC LCA & PROC LTA user's guide version 1.2.5*. University Park, PA: The Methodology Center, The Pennsylvania State University.
- Lanza, S. T., Rhoades, B. L., Greenberg, M. T., Cox, M., & The Family Life Project Key Investigators. (2011). Modeling multiple risks during infancy to predict quality of the caregiving environment: Contributions of a person-centered approach. *Infant Behavior and Development*, 34, 390–406.
- Lanza, S. T., Savage, J. S., & Birch, L. L. (2010). Identification and prediction of latent classes of weight-loss strategies among women. *Obesity*, 18, 833–40.
- Leeb, H., & Pötscher, B. M. (2005). Model selection and inference: Facts and fiction. *Econometric Theory*, 21, 21–59.
- Lenth, R. V. (2007). *Post hoc power: Tables and commentary* (Tech. Rep.). Iowa City, IA: University of Iowa, Department of Statistics and Actuarial Science.
- Lin, T. H., & Dayton, C. M. (1997). Model selection information criteria for non-nested latent class models. *Journal of Educational and Behavioral Statistics*, 22, 249–264.
- Lindsay, B. G. (1995). *Mixture models: Theory, geometry, and applications*. Hayward, CA: Institute for Mathematical Statistics.
- MacCallum, R. C., Browne, M. W., & Sugawara, H. M. (1996). Power analysis and determination of sample size for covariance structure modeling. *Psychological Methods*, 1, 130–149.
- MacCallum, R. C., Lee, T., & Browne, M. W. (2010). The issue of isopower in power analysis for tests of structural equation models. *Structural Equation Modeling*, 17, 23–41.
- MacCallum, R. C., Widaman, K. F., Zhang, S., & Hong, S. (1999). Sample size in factor analysis. *Psychological Methods*, 4, 84–99.
- McLachlan, G. J. (1987). On bootstrapping the likelihood ratio test statistic for the number of components in a normal mixture. *Annals of Statistics*, 36, 318–324.
- McLachlan, G., & Peel, D. (2000). *Finite mixture models*. New York, NY: Wiley.
- Muthén, B. O. (2004). *Mplus technical appendices*. Los Angeles, CA: Muthén & Muthén.
- Muthén, B. (2008). Latent variable hybrids: Overview of old and new models. In G. R. Hancock, & K. M. Samuelsen (Eds.), *Advances in latent variable mixture models* (pp. 1–24). Charlotte, NC: Information Age.
- Muthén, L. K., & Muthén, B. O. (2007). *Mplus users guide* (5th ed.). Los Angeles, CA: Muthén & Muthén.
- Nylund, K. L., Asparouhov, T., & Muthén, B. O. (2007). Deciding on the number of classes in latent class analysis and growth mixture modeling: A Monte Carlo simulation study. *Structural Equation Modeling*, 14, 535–569.
- Preacher, K. J., & MacCallum, R. C. (2002). Exploratory factor analysis in behavior genetics research: Factor recovery with small sample sizes. *Behavior Genetics*, 32, 153–161.
- Raftery, A. E. (1995). Bayesian model selection in social research (with discussion). *Sociological Methodology*, 25, 111–196.
- Ramaswamy, V., DeSarbo, W. S., Reibstein, D. J., & Robinson, W. T. (1993). An empirical pooling approach for estimating marketing mix elasticities with PIMS data. *Marketing Science*, 12, 103–124.
- R Development Core Team. (2010). *R: A language and environment for statistical computing*. Vienna, Austria: R Foundation for Statistical Computing. Retrieved from <http://www.R-project.org>
- Rindskopf, D., & Rindskopf, W. (1986). The value of latent class analysis in medical diagnosis. *Statistics in Medicine*, 5, 21–27.
- Saris, W. E., & Satorra, A. (1993). Power evaluations in structural equation models. In K. A. Bollen & J. S. Long (Eds.), *Testing structural equation models* (pp. 181–204). Newbury Park, CA: Sage.
- Satorra, A., & Saris, W. E. (1985). Power of the likelihood ratio test in covariance structure analysis. *Psychometrika*, 50, 83–90.
- Schwarz, G. (1978). Estimating the dimension of a model. *Annals of Statistics*, 6, 461–464.
- Sclove, S. L. (1987). Application of model-selection criteria to some problems in multivariate analysis. *Psychometrika*, 52, 333–343.
- Shannon, C. E. (1948). A mathematical theory of communication. *The Bell System Technical Journal*, 27, 379–423, 623–656. Retrieved from <http://cm.bell-labs.com/cm/ms/what/shannonday/shannon1948.pdf>
- Shevlin, M., Murphy, J., Dorahy, M. J., & Adamson, G. (2007). The distribution of positive psychosis-like symptoms in the population: A latent class analysis of the National Comorbidity Survey. *Schizophrenia Research*, 89, 101–109.
- Teräsvirta, T., & Mellin, I. (1986). Model selection criteria and model selection tests in regression models. *Scandinavian Journal of Statistics*, 13, 159–171.
- Udry, J. R. (2003). *The National Longitudinal Study of Adolescent Health (Add Health), Waves I & II, 1994–1996; Wave III, 2001–2002* [Machine-readable data file and documentation]. Chapel Hill, NC: Carolina Population Center, University of North Carolina at Chapel Hill.
- Uebersax, J. S., & Grove, W. M. (1990). Latent class analysis of diagnostic agreement. *Statistics in Medicine*, 9, 559–572.
- van der Vaart, A. W. (1998). *Asymptotic statistics*. Cambridge, UK: Cambridge University Press.
- Vermunt, J. K., & Magidson, J. (2004). Local independence. In M. Lewis-Beck, A. Bryman, & T. F. Liao (Eds.), *The Sage encyclopedia of social sciences research methods* (pp. 580–581). Thousand Oaks, CA: Sage. Retrieved from <http://arno.uvt.nl/show.cgi?fid=13296>
- Vermunt, J. K., & Magidson, J. (2005). *Technical guide for Latent GOLD 4.0: Basic and advanced*. Belmont, MA: Statistical Innovations.
- Wahba, G. (1990). *Spline models of observational data*. Philadelphia, PA: SIAM.
- Walters, G. D. (2011). The latent structure of life-course-persistent antisocial behavior: Is Moffitt's developmental taxonomy a true taxonomy? *Journal of Consulting and Clinical Psychology*, 79, 96–105.
- Wood, S. N. (2003). Thin plate regression splines. *Journal of the Royal Statistical Society, Series B*, 65, 95–114.
- Wood, S. N. (2011). Fast stable restricted maximum likelihood and marginal likelihood estimation of semiparametric generalized linear models. *Journal of the Royal Statistical Society, Series B*, 73, 3–36.
- Wood, S. N. (2012). Package mgcv. Retrieved from <http://cran.r-project.org/web/packages/mgcv/mgcv.pdf>
- Wu, Q. (2009). *Class extraction and classification accuracy in latent class models* (Unpublished doctoral dissertation). State College, PA: Pennsylvania State University.
- Xian, H., Scherrer, J. F., Madden P. A., Lyons M. J., Tsuang M., True W. R., & Eisen S. A. (2005). Latent class typology of nicotine withdrawal: Genetic contributions and association with failed smoking cessation and psychiatric disorders. *Psychological Medicine*, 35, 409–419.
- Yang, C. (2006). Evaluating latent class analysis in qualitative phenotype identification. *Computational Statistics & Data Analysis*, 50, 1090–1104.
- Yuan, K.-H., & Hayashi, K. (2003). Bootstrap approach to inference and power analysis based on three statistics for covariance structure models. *British Journal of Mathematical and Statistical Psychology*, 56, 93–110.
- Zucchini, W. (2000). An introduction to model selection. *Journal of Mathematical Psychology*, 44, 41–61.

APPENDIX A CALCULATING THE EFFECT SIZE MEASURES

Equations 3 and 4, for the w and KL effect size measures, suggest that to calculate effect size one needs to calculate P_0 and P_1 values for each of the cells in the contingency table of possible responses (of which there are 2^J if J dichotomous items are being analyzed). Fortunately, this is not as burdensome as it sounds. All that is required are proposed theoretical values for the class sizes γ and response probabilities ρ for H_1 . All of the P_0 and P_1 can then be calculated as outlined here. For dichotomous items, this can be done automatically using a SAS macro provided at <http://methodology.psu.edu/LcaEffectSizes/>, but we provide an outline and the rationale for the procedure here.

1. Take a set of proposed γ values and a set of proposed ρ values provided by the investigator. These values specify the K -class H_1 model that is assumed to be true for calculating power.
2. Given these H_1 parameters, calculate P_1 for each of the cells using Equation 1.
3. Create a synthetic data set having size approximately P_1 for each cell. We do this by first creating a data set with one line for each cell (e.g., [1, 1, 1, 1, 1], [1, 1, 1, 1, 2], etc.) and then giving each cell a frequency weight proportional to the P_1 for that cell. This data set now represents the population probability distribution specified by H_1 , and could be considered a "pseudo-population."
4. Do a $(K - 1)$ -class LCA on the pseudo-population (even though the pseudo-population was generated under a K -class model). Do this LCA very carefully (with many random start values) to try to find the global maximum likelihood solution. Intuitively, the γ and ρ parameters from this solution now provide a fitted $(K - 1)$ -class model that approximates the fit of the true K -class model about as well

as any $(K - 1)$ -class model can do. This in turn approximates the bootstrapped distribution used by the BLRT.

5. Using the calculated best H_0 parameters, calculate P_0 for each of the cells by again using Equation 1.
6. Now that P_0 and P_1 values are available for all cells, the effect size can be calculated directly using Formulas 3 and 4.

In the preceding outline, one might ask why we calculate the H_0 probabilities indirectly from the H_1 probabilities, instead of asking the investigator for γ and ρ values for H_0 as well. In other words, why are H_0 and H_1 being treated so differently, and why is the elaborate step of constructing a pseudo-population necessary? The answer is that this corresponds to the way the BLRT test is done in practice. When doing a BLRT in practice, the bootstrap H_0 data sets are simulated not from a prespecified set of parameters, but from the best fitting $(K - 1)$ -class parameters to the observed data. Thus, to be realistic, after we propose a particular assumed set of H_1 parameters for calculating P_1 , we should not also propose another arbitrary set of H_0 parameters for calculating P_0 . Instead, we need to find the set of H_0 parameters that give P_0 values that are, in general, as close as possible to P_1 . This sounds difficult but is easy: As described earlier, one can just use existing LCA software to fit a $(K - 1)$ -class model to a sample of fake data having cell proportions equal to the P_1 . Because these steps are automated, the process is actually much simpler for the user than having to specify separate theoretical values for both the H_0 model and the H_1 model.

APPENDIX B

Appendix B is provided online at <http://methodology.psu.edu/media/LcaEffectSizes/AppendixDziakLanzaTan.pdf>